



Detailed Report


Model Config: yolov8s_config_0



Latency Breakdown for Online Performance of yolov8s_config_0




GPU Memory vs. Latency curves for config yolov8s_config_0



GPU Utilization vs. Latency curves for config yolov8s_config_0

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
8	2272.498	2038.12	1515.322	0.031	515.855	3.81099	0	0
4	1209.948	922.628	454.0	0.027	461.91	4.26421	0	0
2	605.332	527.443	0.162	0.025	520.739	3.73816	0	0
1	232.089	211.643	0.065	0.023	208.176	4.64098	0	0



GPU Power vs. Latency curves for config yolov8s_config_0

The model config **yolov8s_config_0** uses 2 CPU instances with a max batch size of 1 and has dynamic batching enabled. 4 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.