



Detailed Report


Model Config: yolov8s_config_2



Latency Breakdown for Online Performance of yolov8s_config_2




GPU Memory vs. Latency curves for config yolov8s_config_2



GPU Utilization vs. Latency curves for config yolov8s_config_2

GPU Power vs. Latency



GPU Power vs. Latency curves for config yolov8s_config_2

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	24937.74	23138.278	21613.545	0.027	1447.701	5.50464	0	0
8	13540.525	11799.463	10216.883	0.027	1505.884	5.28703	0	0
4	6172.181	5891.329	4353.549	0.026	1469.326	5.42182	0	0
2	4609.052	2983.714	1417.175	0.026	1494.783	5.30959	0	0
16	3933.078	2739.747	1527.861	19.422	1164.983	5.84847	0	0
1	1740.333	1471.762	0.07	0.025	1441.944	5.40573	0	0

The model config **yolov8s_config_2** uses 1 CPU instance with a max batch size of 8 and has dynamic batching enabled. 6 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.