



Detailed Report


Model Config: yolov8s_config_3




Latency Breakdown for Online Performance of yolov8s_config_3



GPU Memory vs. Latency curves for config yolov8s_config_3



GPU Utilization vs. Latency curves for config yolov8s_config_3



GPU Power vs. Latency curves for config yolov8s_config_3

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
64	27630.118	26299.696	19676.661	81.672	6463.796	4.86322	0	0
64	13258.816	13109.665	6541.132	103.857	6394.339	4.85038	0	0
8	5528.994	3618.759	1016.096	20.701	2546.461	4.38717	0	0
4	2929.087	1955.654	553.254	2.883	1380.917	4.04922	0	0
2	1544.241	1045.634	1.026	0.054	1034.49	3.79452	0	0
1	464.305	399.001	0.813	0.025	391.274	4.94731	0	0

The model config **yolov8s_config_3** uses 2 CPU instances with a max batch size of 16 and has dynamic batching enabled. 6 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.