



Detailed Report


Model Config: yolov8s_config_4




Latency Breakdown for Online Performance of yolov8s_config_4



GPU Memory vs. Latency curves for config yolov8s_config_4



GPU Utilization vs. Latency curves for config yolov8s_config_4



GPU Power vs. Latency curves for config yolov8s_config_4

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
24	11967.441	11604.536	8679.637	13.955	2879.073	4.1192	0	0
8	4925.597	3892.456	950.197	13.642	2900.58	4.00805	0	0
4	3244.919	2168.742	420.569	0.602	1732.183	3.65271	0	0
2	1256.622	1076.106	1.09	0.058	1065.419	3.69174	0	0
1	602.488	400.497	1.18	0.025	392.54	4.93098	0	0

The model config **yolov8s_config_4** uses 3 CPU instances with a max batch size of 4 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.