



Detailed Report


Model Config: yolov8s_config_5




Latency Breakdown for Online Performance of yolov8s_config_5



GPU Memory vs. Latency curves for config yolov8s_config_5



GPU Utilization vs. Latency curves for config yolov8s_config_5



GPU Power vs. Latency curves for config yolov8s_config_5

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	4429.737	3740.95	1826.386	15.37	1872.61	4.09383	0	0
8	2501.467	1947.227	447.565	7.774	1472.863	4.05604	0	0
4	1575.016	1062.756	283.739	2.235	767.467	3.71344	0	0
2	713.517	551.768	0.146	0.042	545.082	3.56288	0	0
1	284.642	233.463	0.077	0.025	229.803	4.1981	0	0

The model config **yolov8s_config_5** uses 2 CPU instances with a max batch size of 4 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.