



Detailed Report


Model Config: yolov8s_config_6




Latency Breakdown for Online Performance of yolov8s_config_6



GPU Memory vs. Latency curves for config yolov8s_config_6



GPU Utilization vs. Latency curves for config yolov8s_config_6



GPU Power vs. Latency curves for config yolov8s_config_6

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
64	14127.823	13207.595	6595.648	107.154	6419.175	4.90297	0	0
8	2963.732	1984.706	363.636	10.04	1592.065	3.93873	0	0
4	1583.383	1062.062	292.333	4.228	755.962	3.72158	0	0
2	692.59	551.828	1.131	0.027	544.7	3.58278	0	0
1	263.634	233.649	1.129	0.025	228.698	4.19642	0	0

The model config **yolov8s_config_6** uses 2 CPU instances with a max batch size of 16 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.