



Detailed Report


Model Config: yolov8s_config_7




Latency Breakdown for Online Performance of yolov8s_config_7



GPU Memory vs. Latency curves for config yolov8s_config_7



GPU Utilization vs. Latency curves for config yolov8s_config_7



GPU Power vs. Latency curves for config yolov8s_config_7

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
32	24330.799	24186.525	21135.212	50.642	2943.424	5.31682	0	0
8	8487.778	6071.803	3426.629	39.775	2550.471	5.15068	0	0
4	3926.406	3055.349	1070.992	14.164	1918.114	5.18668	0	0
2	1660.156	1446.353	682.819	0.08	731.572	5.47287	0	0
1	791.778	722.796	0.793	0.025	706.497	5.48535	0	0

The model config **yolov8s_config_7** uses 1 CPU instance with a max batch size of 16 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.