



Detailed Report


Model Config: yolov8s_config_10




Latency Breakdown for Online Performance of yolov8s_config_10



GPU Memory vs. Latency curves for config yolov8s_config_10



GPU Utilization vs. Latency curves for config yolov8s_config_10



GPU Power vs. Latency curves for config yolov8s_config_10

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
8	3620.738	2498.736	1542.473	0.029	946.392	3.12385	0	0
6	2012.382	1843.679	907.171	0.032	927.269	3.19566	0	0
4	1782.521	1232.555	303.189	0.027	920.566	3.24171	0	0
2	867.359	631.466	0.198	0.025	625.056	3.16067	0	0
1	400.68	235.954	0.071	0.024	232.57	4.16021	0	0

The model config **yolov8s_config_10** uses 3 CPU instances with a max batch size of 1 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.