



Detailed Report


Model Config: yolov8s_config_11



Latency Breakdown for Online Performance of yolov8s_config_11




GPU Memory vs. Latency curves for config yolov8s_config_11



GPU Utilization vs. Latency curves for config yolov8s_config_11

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
8	1359.65	1186.102	1031.458	0.035	149.024	6.55354	0	0
4	757.514	590.957	438.2	0.064	146.862	6.63006	0	0
2	480.286	330.996	161.89	0.029	163.37	5.93879	0	0
1	192.988	145.582	0.073	0.023	142.298	6.71599	0	0



GPU Power vs. Latency curves for config yolov8s_config_11

The model config **yolov8s_config_11** uses 1 CPU instance with a max batch size of 1 and has dynamic batching enabled. 4 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.