



Detailed Report


Model Config: yolov8s_config_12




Latency Breakdown for Online Performance of yolov8s_config_12



GPU Memory vs. Latency curves for config yolov8s_config_12



GPU Utilization vs. Latency curves for config yolov8s_config_12



GPU Power vs. Latency curves for config yolov8s_config_12

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
16	12177.968	11965.779	10432.678	17.5	1471.592	5.36583	0	0
8	6761.917	6027.22	4497.357	23.908	1470.774	5.29939	0	0
4	4771.982	3042.024	1494.93	25.186	1486.746	5.1692	0	0
16	3365.56	2992.715	1439.811	31.325	1487.959	5.04009	0	0
2	3304.727	1479.602	641.916	0.751	805.415	5.36328	0	0
1	744.393	682.227	0.091	0.025	668.603	5.81815	0	0

The model config **yolov8s_config_12** uses 1 CPU instance with a max batch size of 8 and has dynamic batching enabled. 6 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.