



# Detailed Report


Model Config: yolov8s\_config\_13




Latency Breakdown for Online Performance of yolov8s\_config\_13



GPU Memory vs. Latency curves for config yolov8s\_config\_13



GPU Utilization vs. Latency curves for config yolov8s\_config\_13



GPU Power vs. Latency curves for config yolov8s\_config\_13

| Request Concurrency | p99 Latency (ms) | Client Response Wait (ms) | Server Queue (ms) | Server Compute Input (ms) | Server Compute Infer (ms) | Throughput (infer/sec) | Max GPU Memory Usage (MB) | Average GPU Utilization (%) |
|---------------------|------------------|---------------------------|-------------------|---------------------------|---------------------------|------------------------|---------------------------|-----------------------------|
| 8                   | 3355.951         | 2903.724                  | 2143.181          | 6.214                     | 732.411                   | 5.32037                | 0                         | 0                           |
| 4                   | 1576.741         | 1392.34                   | 678.639           | 5.005                     | 688.17                    | 5.67798                | 0                         | 0                           |
| 8                   | 1534.657         | 1339.653                  | 644.599           | 7.984                     | 663.124                   | 5.72326                | 0                         | 0                           |
| 2                   | 960.233          | 638.047                   | 295.865           | 0.137                     | 329.818                   | 6.19407                | 0                         | 0                           |
| 1                   | 351.491          | 311.291                   | 0.074             | 0.022                     | 304.673                   | 6.31655                | 0                         | 0                           |

The model config **yolov8s\_config\_13** uses 1 CPU instance with a max batch size of 4 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.