



Detailed Report


Model Config: yolov8s_config_default



Latency Breakdown for Online Performance of yolov8s_config_default




GPU Memory vs. Latency curves for config yolov8s_config_default



GPU Utilization vs. Latency curves for config yolov8s_config_default

GPU Power vs. Latency



GPU Power vs. Latency curves for config yolov8s_config_default

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
32	5200.537	5032.135	2467.095	65.316	2447.583	6.19595	0	0
8	1742.976	1361.677	532.139	9.188	801.153	5.76699	0	0
4	995.086	644.624	213.167	2.02	417.965	6.10643	0	0
2	464.174	295.964	128.469	0.243	161.695	6.64274	0	0
1	171.346	145.071	0.072	0.023	141.697	6.72288	0	0

The model config **yolov8s_config_default** uses 1 CPU instance with a max batch size of 16 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) with total memory 0 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.